

# Reducing memory allocations

–

## The case of C++ associative containers

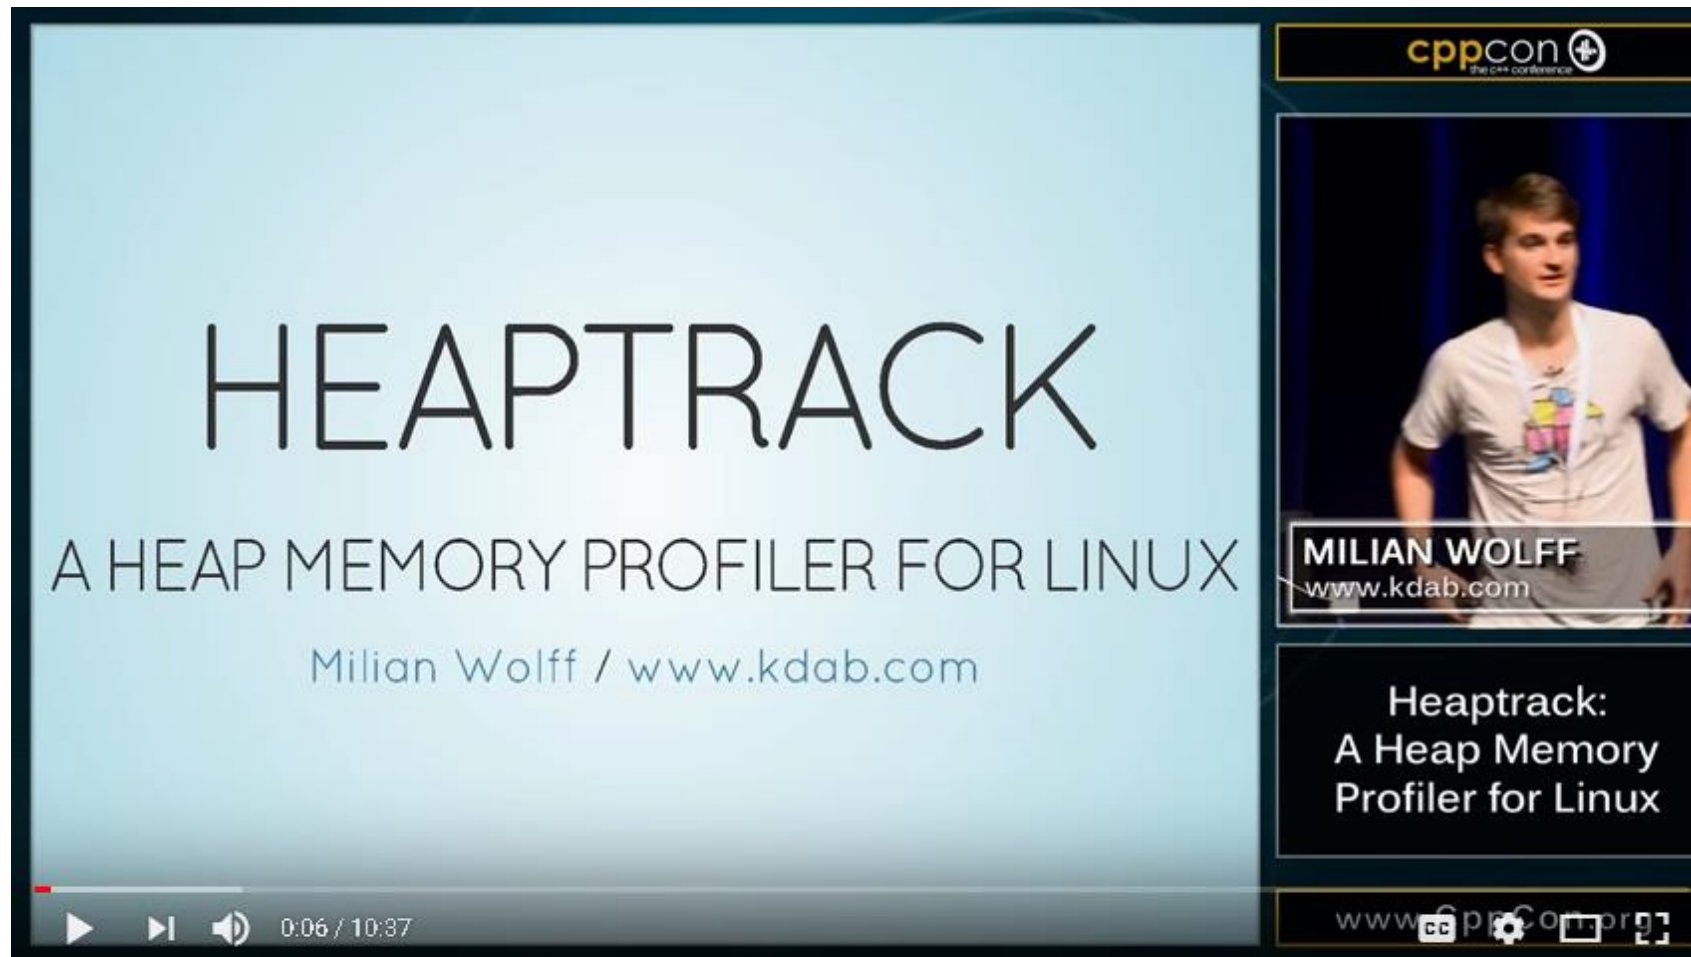
Arnaud Desitter

ACCU Oxford

7 June 2018

# Agenda

- Measuring memory allocations
- Memory allocations of associative containers
- How to improve them



CppCon 2015: Milian Wolff "Heaptrack: A Heap Memory Profiler for Linux"

**Tip:** heaptrack\_gui is easy to build on a recent version of Linux (e.g. Ubuntu 18)

Source: youtube.com

# C++ associative containers

	<code>std::map</code> <code>std::set</code>	<code>std::unordered_map</code> <code>std::unordered_set</code>
	Tree based	Hash based
Lookup complexity	$O(\log(n))$	$O(1)$ amortised, $O(n)$ worst case
Functional requirement	A comparison function	A hash and an equality function
Source of dynamic allocations	Nodes	(Somewhat slimmer) nodes and buckets

# std::set

```
#include <set>

void ex1() {
    std::set<std::size_t> set;
    constexpr std::size_t nmax = 1000;
    for (std::size_t i = 0; i != nmax; ++i)
        set.emplace(345 + i);
}

int main() { ex1(); }
```

```
>heaptrack ./a.out
heaptrack stats:
      allocations:           1000
```

Allocations ^	Temporar	Peak	Leaked	Allocated	Location
> 1000	1	40.0 kB	0 B	40.0 kB	__gnu_cxx::new_allocator<std::_Rb_tree_node<unsigned long> >::allocate(unsigned

# std::unordered\_set (1)

```
#include <unordered_set>

void ex1() {
    std::unordered_set<std::size_t> set;
    constexpr std::size_t nmax = 1000;
    for (std::size_t i = 0; i != nmax; ++i)
        set.emplace(345 + i);
}

int main() { ex1(); }
```

```
>heaptrack ./a.out
heaptrack stats:
      allocations:           1009
```

Allocations ^	Temporar	Peak	Leaked	Allocated	Location
> 1000	1	11.3 kB	0 B	16.0 kB	__gnu_cxx::new_allocator<std::__detail::_Hash_node<unsigned long, false> >::alloc
> 9	0	17.6 kB	0 B	22.8 kB	__gnu_cxx::new_allocator<std::__detail::_Hash_node_base*>::allocate(unsigned l

Slimmer nodes

Buckets

# std::unordered\_set (2)

```
#include <unordered_set>

void ex1() {
    std::unordered_set<std::size_t> set;
    constexpr std::size_t nmax = 1000;
    set.reserve(nmax);
    for (std::size_t i = 0; i != nmax; ++i)
        set.emplace(345 + i);
}

int main() { ex1(); }
```


```
>heaptrack ./a.out
heaptrack stats:
      allocations:           1001
```

Allocations ^	Temporar	Peak	Leaked	Allocated	Location
> 1000	1	16.0 kB	0 B	16.0 kB	__gnu_cxx::new_allocator<std::__detail::Hash_node<unsigned long, false> >::alloca
> 1	0	8.2 kB	0 B	8.2 kB	__gnu_cxx::new_allocator<std::__detail::Hash_node_base*>::allocate(unsigned long

Associative containers will call the allocator at least “size” times.




# Local allocators to the rescue!



Local (arena) Memory Allocators - John Lakos [ACCU 2017]

cppcon | 2017  
THE C++ CONFERENCE • BELLEVUE, WASHINGTON



PABLO HALPERN

## Allocators: The Good Parts

### Allocators got a whole lot easier

Task	C++98/C++03	C++11/C++14	C++17 polymorphic_allocator<byte>
Use an allocator	MEDIUM viral templates	MEDIUM viral templates	EASY
Create an allocator	MEDIUM Lots of boilerplate, non-portable state	EASY	EASY just derive from memory_resource
Create a scoped allocator	IMPOSSIBLE	MEDIUM-EASY alias scoped_allocator_adaptor	EASY polymorphic_allocator is scoped
Create a new allocator-aware container	MEDIUM rebinding needed, ignore allocator state?	HARD propagation traits, allocator_traits	EASY skip C++11 complexity

Pablo Halpern, 2017 (CC BY 4.0) 9/29/2017 21

22:31 / 1:00:48

CC HD

CppCon 2017: Pablo Halpern "Allocators: The Good Parts"

By the author of the C++17 allocators spec N3916.  
The first 30 minutes are the most important.

## std::map

“Allocator” is a customisation point.  
By default, new and delete are used.

Defined in header <map>

```
template<
    class Key,
    class T,
    class Compare = std::less<Key>,
    class Allocator = std::allocator<std::pair<const Key, T>> >
> class map;                                     (1)
```

```
namespace pmr {
    template<class Key, class T, class Compare = std::less<Key>>
    using map = std::map<Key, T, Compare,
                        std::pmr::polymorphic_allocator<std::pair<const Key, T>>> >
}                                                (2) (since C++17)
```

“pmr” stands for “polymorphic memory resource”.

# What is the point of allocators and resources?

- Coalesce allocations of small blocks in pre-allocated larger blocks
  - Better locality
  - Less heap fragmentation
- Available memory resources in **C++17**:
  - `synchronized_pool_resource`: thread-safe pools of similar-sized memory pools
  - `unsynchronized_pool_resource`: non-thread-safe version
  - **`monotonic_buffer_resource`**: *very fast*, non-thread-safe allocation into buffers with do-nothing deallocation
    - Usually what we need for most use of associative containers if no erasure take place.

# std::set

```
#include <set>

void ex1() {
    std::set<std::size_t> set;
    constexpr std::size_t nmax = 1000;
    for (std::size_t i = 0; i != nmax; ++i)
        set.emplace(345 + i);
}

int main() { ex1(); }
```

```
>heaptrack ./a.out
heaptrack stats:
    allocations:           1000
```

Allocations	Temporar	Peak	Leaked	Allocated	Location
> 1000	1	40.0 kB	0 B	40.0 kB	__gnu_cxx::new_allocator<std::_Rb_tree_node<unsigned long>>::allocate(unsigned

# std::pmr::set

```
#include <memory_resource>
#include <set>
void ex1() {
    std::pmr::monotonic_buffer_resource res;
    std::pmr::set<std::size_t> map{ &res };
    constexpr std::size_t nmax = 1000;
    for (std::size_t i = 0; i != nmax; ++i)
        map.insert(345 + i);
}
int main() { ex1(); }
```

Unfortunately, no recent  
compilers implement the “pmr”  
facilities.  
(with the notable exception of Visual Studio®  
2017 update 6)

# boost::container::pmr::set

```
#include <boost/container/pmr/monotonic_buffer_resource.hpp>
#include <boost/container/pmr/set.hpp>

void ex1() {
    namespace pmr = boost::container::pmr;
    pmr::monotonic_buffer_resource res;
    pmr::set<std::size_t> map{ &res };
    constexpr std::size_t nmax = 1000;
    for (std::size_t i = 0; i != nmax; ++i)
        map.insert(345 + i);
}

int main() { ex1(); }
```

```
>heaptrack ./a.out
heaptrack stats:
      allocations:           7
```

Allocations ^	Temporar	Peak	Leaked	Allocated	Location
> 7	1	32.6 kB	0 B	32.6 kB	boost::container::pmr::monotonic_buffer_resource::do_allocate(unsig

# Alternatives

- Use an open source implementation
  - Boost (limited to boost containers)
  - <https://github.com/mmcshane/pmr>
  - <https://github.com/phalpern/CppCon2017Code>
- Roll your own:
  1. A subset of `std::pmr::monotonic_buffer_resource`,
  2. A simple, usable allocator optimized for node based containers,
  3. A series of type aliases



# Home-brewed implementation (1)

// In C++17, when available, we can replace the code below with `std::pmr::monotonic_buffer_resource`.

```
namespace project { namespace pmr {  
    class MonotonicBufferResource {  
    public:  
        MonotonicBufferResource();  
        ~MonotonicBufferResource();  
  
        void* allocate(std::size_t num_bytes, std::size_t alignment);  
        void deallocate(void*)  
        {} // Do not de-allocate individual blocks.  
    private:  
        struct PImpl; // see for instance https://github.com/mmcshane/pmr  
        std::unique_ptr<PImpl> m_impl;  
    };  
}}
```

# Home-brewed implementation (2)

```
template<typename T>
class NodeAllocator {
    MonotonicBufferResource* m_res;
public:
    using value_type = T;    // Refer to https://howardhinnant.github.io/allocator\_boilerplate.html
    value_type* allocate(std::size_t numObjects) {
        if (numObjects != 1)
            return static_cast<value_type*>(::operator new(sizeof(value_type)*numObjects));
        else
            return static_cast<value_type*>(m_res->allocate(sizeof(value_type), alignof(value_type)));
    }
    void deallocate(void* p, std::size_t numObjects) {
        if (numObjects != 1)
            ::operator delete(p);
        else
            m_res->deallocate(p);
    }
};
```

# Home-brewed implementation (3)

// In C++17, when available, we can replace the code below with `std::pmr::set`.

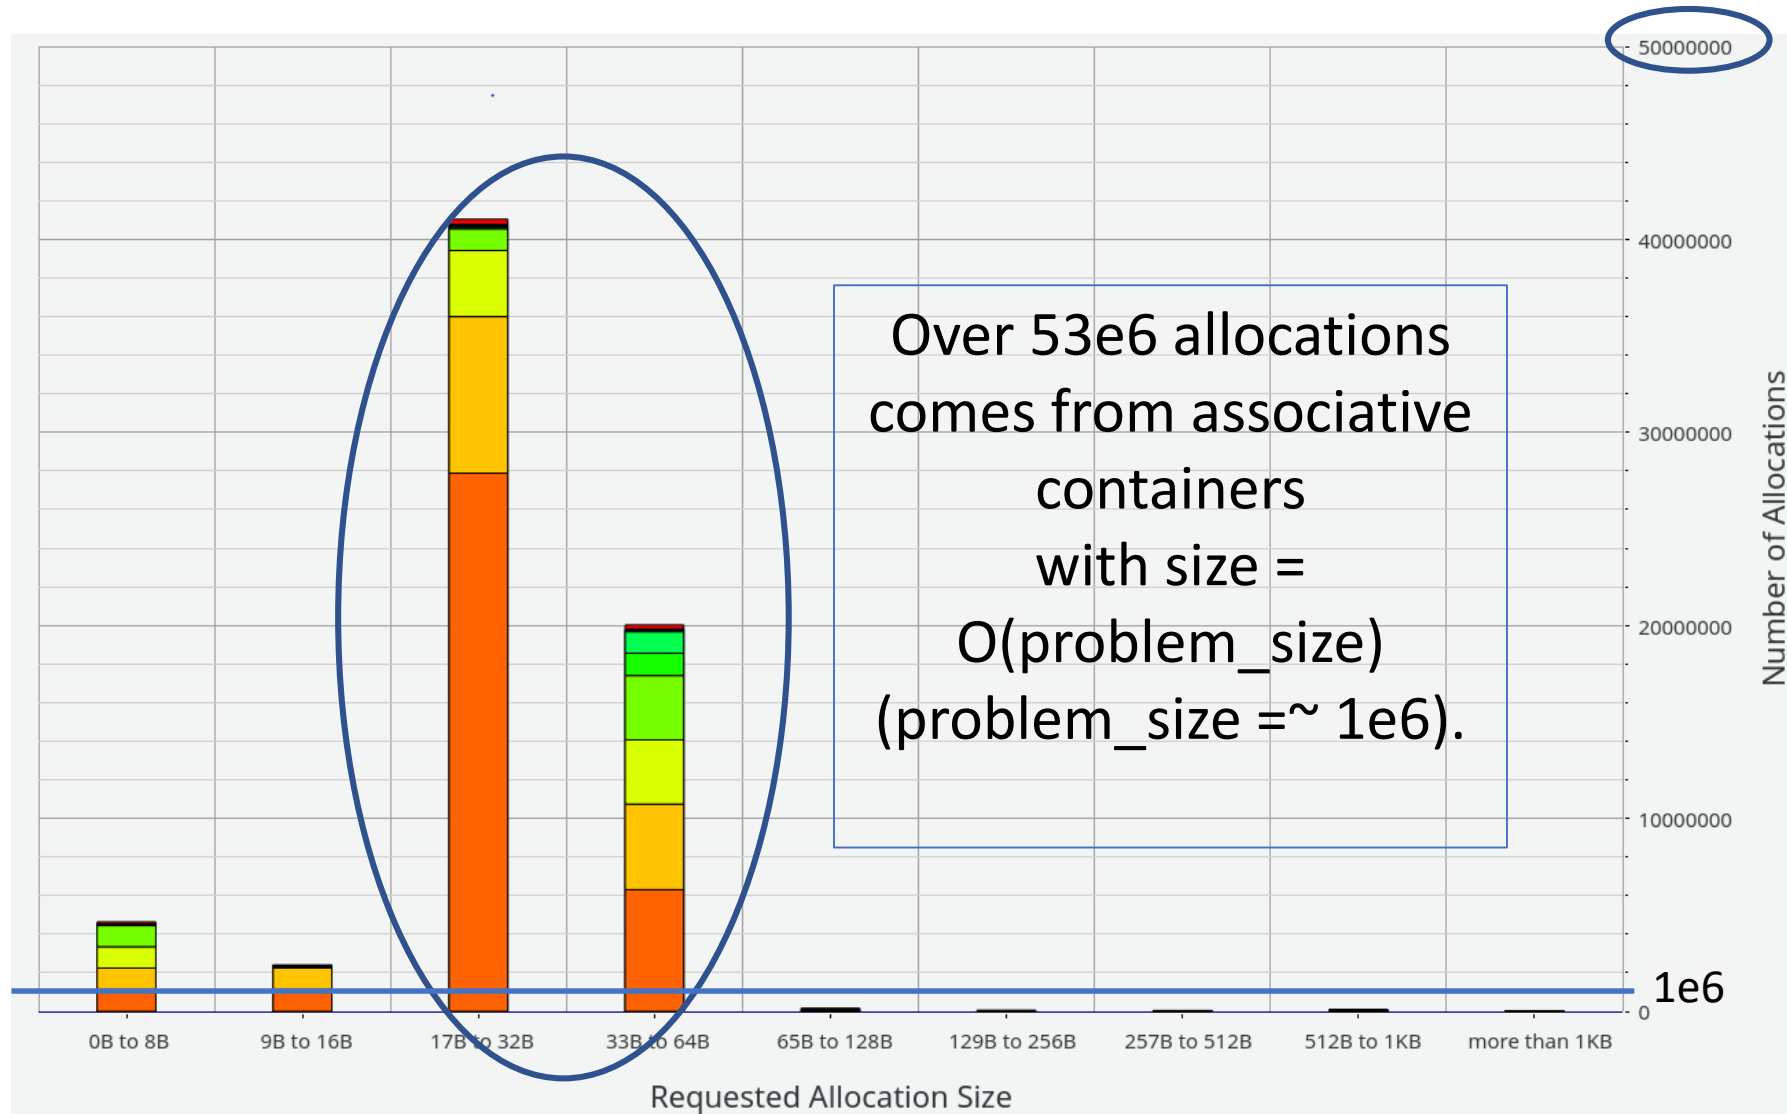
```
namespace project { namespace pmr {  
    template <typename Key,  
              typename Compare = std::less<Key>>  
    using set = std::set<Key, Compare,  
                        project::pmr::NodeAllocator<Key>>;  
  
    template <typename Key,  
              typename Hash = std::hash<Key>,  
              typename Pred = std::equal_to<Key>>  
    using unordered_set = std::unordered_set<Key, Hash, Pred,  
                                              project::pmr::NodeAllocator<Key>>;  
}}
```

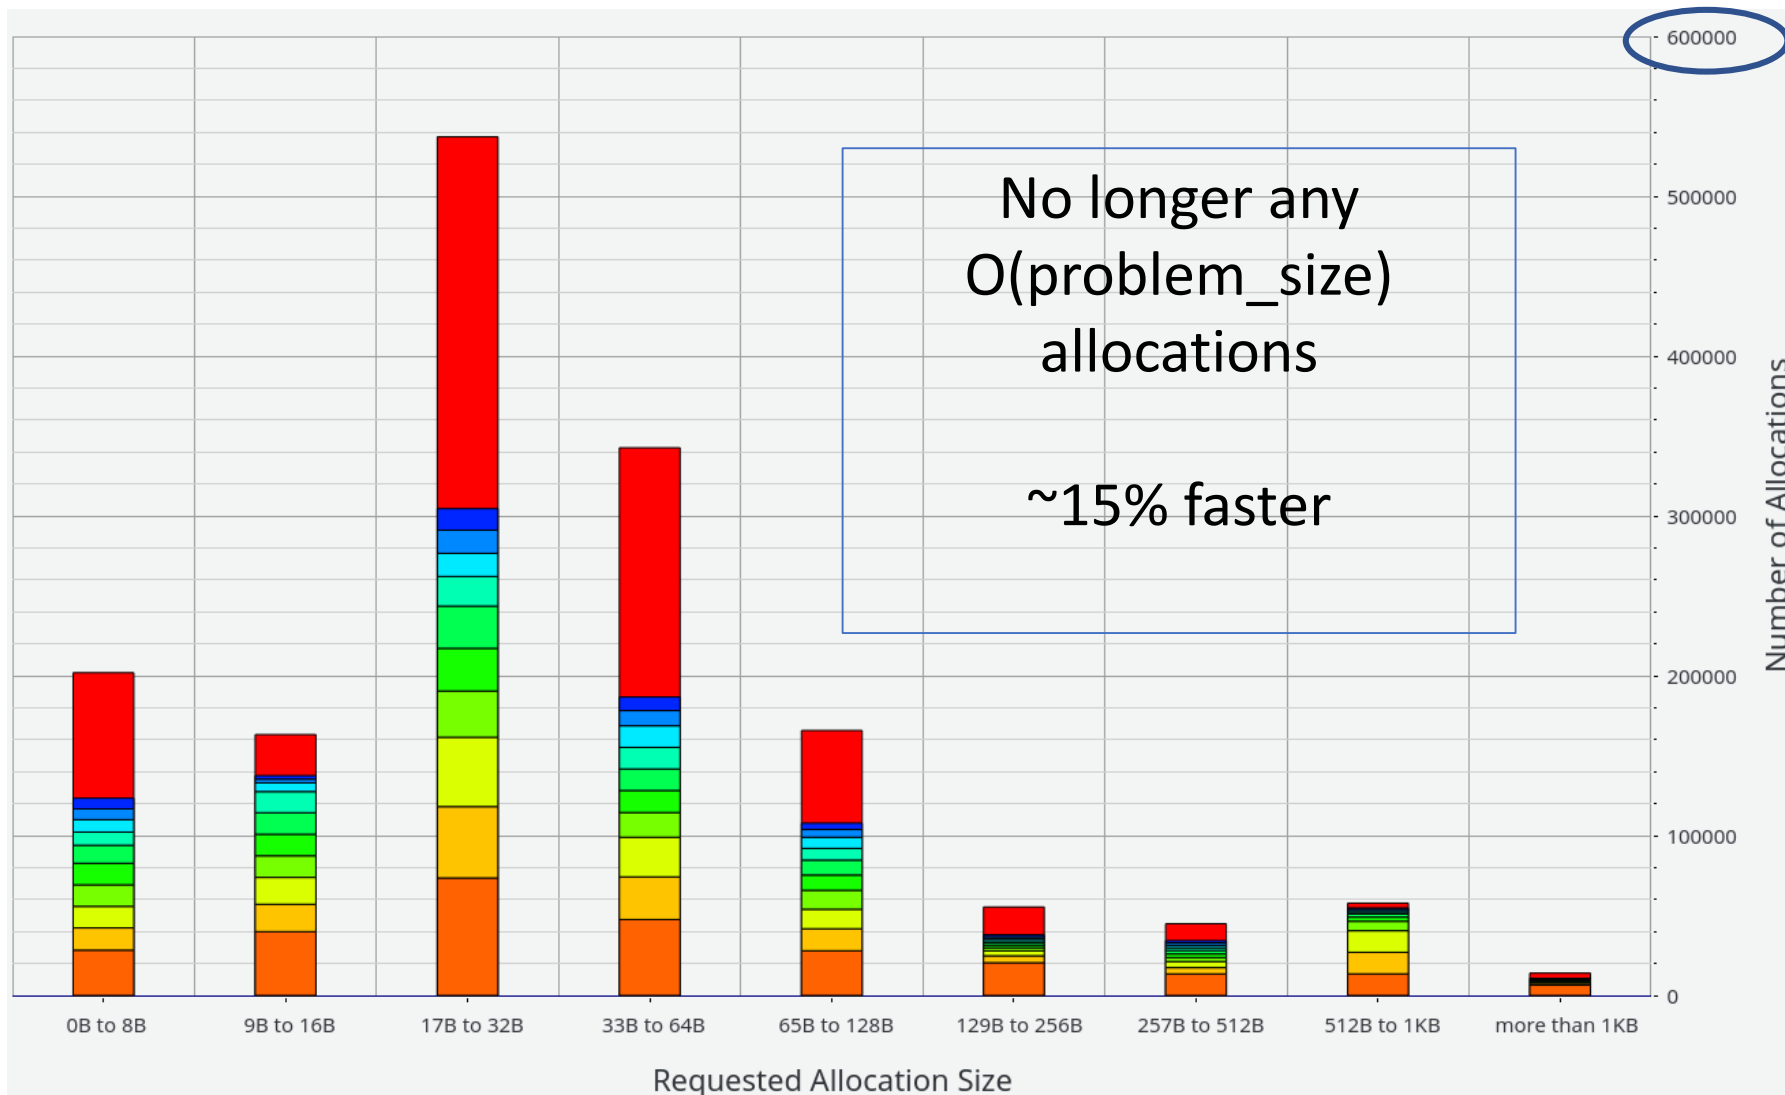
# Usage

```
std::unordered_set<std::int64_t> vertex_ids;
```

->

```
project::pmr::MonotonicBufferResource monotonic_res;  
project::pmr::unordered_set<std::int64_t> vertex_ids{&monotonic_res};
```





# Conclusions

- Memory allocations can be surprisingly expensive.
- Heaptrack rocks.
- Use C++17 local memory allocators when appropriate.