

Reducing memory allocations in a large C++ application

Arnaud Desitter

ACCU Oxford


18 July 2019

Roadmap

- Motivations
- Methodology
- Part I: a case study
- Part II: solutions to address excessive memory allocations
 - Vocabulary types
 - Patterns and anti-patterns
- Part III: C++17 pmr allocators (briefly)
- Conclusions

Motivations

Custom allocators are a much discussed topic in the C++ industry.




Local (arena) Memory Allocators - John Lakos [ACCU 2017]

Motivations

Custom allocators are a much discussed topic in the C++ industry.

How do I quantify the memory allocations of my application ?



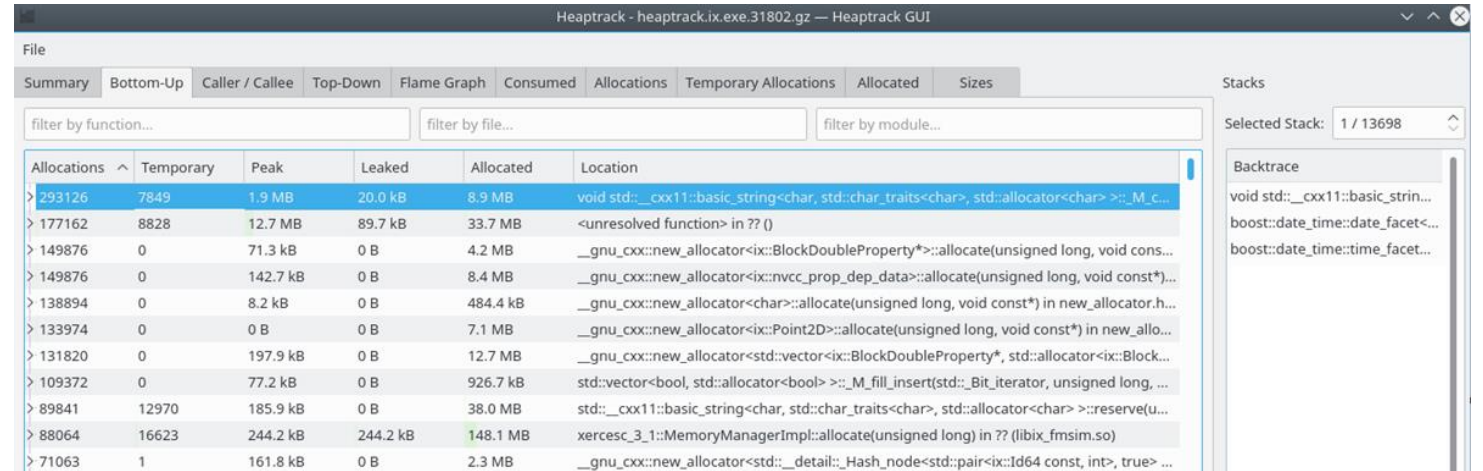
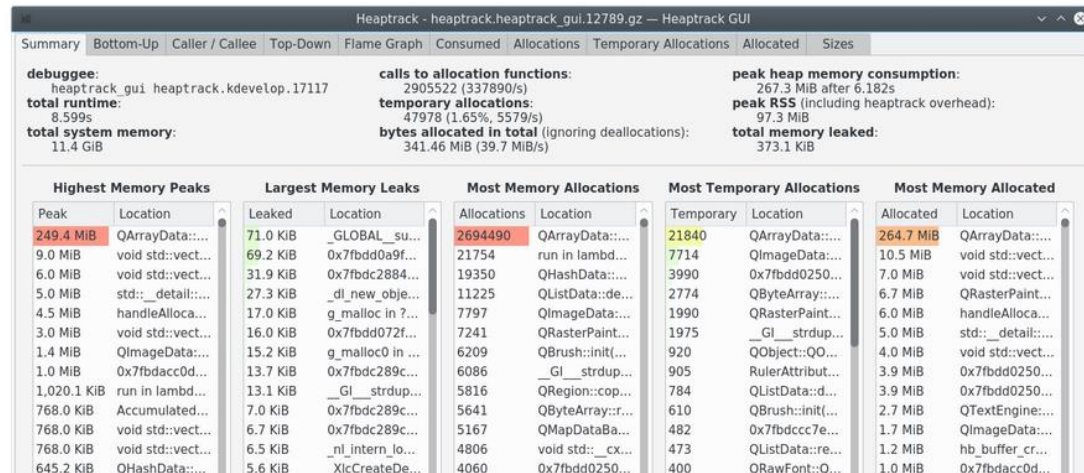
Local (arena) Memory Allocators - John Lakos [ACCU 2017]

Motivations



CppCon 2015: Milian Wolff "Heaptrack: A Heap Memory Profiler for Linux"

Heaptrack




Heaptrack



Flamecharts



Cumulated allocations



Sizes



Consumed

Tips

- Go to conferences !
or watch them on YouTube.
- Do not be afraid to ask questions.
at conferences or on the web.
- Try new tools.
... and make improvements thanks to them.

Milian Wolff's heaptrack

```
# Build it from source (search for "heaptrack build ubuntu") or use pre-built AppImage
# See https://download.kde.org/stable/heaptrack/1.1.0/heaptrack-v1.1.0-x86\_64.AppImage.mirrorlist
# wget .../download.kde.org/stable/heaptrack/1.1.0/heaptrack-v1.1.0-x86_64.AppImage
# chmod +x heaptrack-v1.1.0-x86_64.AppImage
```

```
# heaptrack operates on any executable. Stack traces available only with debugging symbols (compile with "-g")
```

```
> ./heaptrack-v1.1.0-x86_64.AppImage /bin/ls
```

(1) Execution with heaptrack
data collection

```
heaptrack output will be written to "heaptrack.ls.1877.zst"
```

```
starting application, this might take some time...
```

```
heaptrack.ls.1877.zst heaptrack-v1.1.0-x86_64.AppImage
```

```
heaptrack stats:
```

```
    allocations:          44
```

```
    leaked allocations:    38
```

```
    temporary allocations:  2
```

```
Heaptrack finished! Now run the following to investigate the data:
```

```
heaptrack --analyze "heaptrack.ls.1877.zst"
```

```
> ./heaptrack-v1.1.0-x86_64.AppImage --analyze heaptrack.ls.1877.zst
```

(2) Data visualisation

Demo time

Profiling: a methodology


- Choose a case.
- Run it under a profiler.
- Spot a problem.
- Try to fix it.
- Profile again with the fix.
 - Discard fix if it does not work.
 - Submit if it does.
- Iterate until there are no more opportunities for change.
- Choose another case and iterate.

Part I: a case study

A numerical simulation


- Reasonably large synthetic problem
- ... but shortened to a single day of simulation
- Dominated by floating point computation
- Heavily optimised over the years
- Run **without** any concurrency for the sake of this example
 - Considerably faster when many processes and cores are used
 - Allocations costs are amortised across processes.

A case study



Cumulated
allocations


A case study



Before

~630e6

Number of allocations
reduced by x250




~2.5e6

After


A case study

Most allocations were for 8 bytes or less.



Large single source contributors

Before



Largest contributors

After

Cost of allocations are hard to predict

- Looking **carefully** with perf, about **5.5%** time was spent in malloc/free.
- We reduced it to **1.5%**.
- But we have a speed up well above 10%.

Reported time spent in allocator is an **under-estimate** of possible gain as excessive memory allocations may be a source of cache-misses.

Conclusions so far

- Heaptrack works really well.
 - Reasonable information with debug information provided by gcc and clang
 - Overhead is quite small once the number of allocations is under control.
 - Compatible with threads and MPI.
- It helped us to identify and then obtain significant speedup
 - Biggest gains come from removing allocations within loops.
 - Biased towards small size allocations.

Part II: solutions

- Do not copy if you can.
- Avoid allocations.
- Re-use allocated memory.
- Use contiguous containers.

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`gsl::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Part II: solutions

- Do not copy if you can.
 - **Avoid unused objects.**
 - Use references.
 - Use views (`gsl::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Case study – problem #1 – Avoid unnecessary copies

Some leftovers of a debugging experiment: 210e6 allocations

```
std::vector<double> average_m(m.size(), 0.0);  
for (int ic = 0; ic < m.size(); ic++)  
    average_m[ic] = m[ic];
```

->

```
auto const& average_m = m;
```

ALWAYS DO LESS WORK



EFFICIENCY WITH ALGORITHMS, PERFORMANCE WITH DATA STRUCTURES

Chandler Carruth

▶ ▶▶ 32:16 / 1:13:40



Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - **Use references.**
 - Use views (`gsl::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Do not copy – “auto const& x = f()”

```
for (...) { for (...) {  
    std::vector<int> bs = the_sizes[...]->get_sizes();  
->  
for (...) { for (...) {  
    // See GotW #88: A Candidate For the “Most Important const”  
    auto const& bs = the_sizes[...]->get_sizes();
```

Returns “std::vector<int> const&”



Do not copy arguments

// copy !

```
void add_res_from_buffer(std::vector<double> dbuffer);
```

->

```
void add_res_from_buffer(std::vector<double> const& dbuffer);
```

Beware of stateful functors

```
class SortFunctor {  
    std::vector<std::string> m_names; // ...  
};  
std::sort(begin(nodes), end(nodes), SortFunctor{names});
```

Copy!



->

```
class SortFunctor {  
    std::vector<std::string> m_names; // ...  
    SortFunctor(SortFunctor const&) = delete;  
    SortFunctor& operator=(SortFunctor const&) = delete;  
};  
const SortFunctor cmp{names}  
std::sort(begin(nodes), end(nodes), cref(cmp));
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - **Use views (`gsl::span`, `std::string_view`).**
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Vocabulary types - views

- Views are **pointers**.
 - They can dangle.
 - They are cheap to copy.
- C++17 `std::string_view`
 - A non-owning read-only view to a **contiguous** sequence of “char”.
 - May **not** have a trailing ‘\0’.
- GSL `gsl::span<T>` a.k.a. C++20 `std::span<T>`
 - A non-owning view to a **contiguous** sequence of T.
 - `gsl::span<T>`: read-write
 - `gsl::span<const T>`: read-only

Vocabulary types - views

- `std::string_view`

- Avoid creating temporary `std::string` when passing a literal (“Hello world”).
- Frequent solution to excessive allocations due to `std::string`.

- `gsl::span`

- Usually used as function parameters
 - “`std::vector<double> &`” -> `gsl::span<double>` (if `std::vector` is not resized)
 - “`std::vector<double> const&`” -> `gsl::span<const double>`
- Allow passing any contiguous sequences such as:
 - `std::vector`
 - `std::array`
 - `boost::container::small_vector`, `boost::container::static_vector`

Use std::string_view

```
int size = 0;  
const std::unique_ptr<char[]> pchar = block->get_char(size);  
std::string str(pchar.get(), size - 1); // Copy!
```

->

```
int size = 0;  
const std::unique_ptr<char[]> pchar = block->get_char(size);  
const std::string_view str{ pchar.get(), size - 1 };
```

string_view and std::map

```
class Map {  
    std::map<std::string, int> m_map;  
public:  
    bool contain(std::string const& key)  
    { return m_map.find(key) != m_map.end(); }  
};  
Map m; m.contain("a very long string"); // Temporary std::string created
```

->

```
class Map {  
    std::map<std::string, int, std::less<>> m_map; // C++14 transparent comparator  
public:  
    bool contain(std::string_view key) {  
        { return m_map.find(key) != m_map.end(); }  
    };  
Map m; m.contain("a very long string"sv);
```

string_view and std::map

```
class Map {  
    std::map<std::string, int> m_map;  
public:  
    bool contain(std::string const& key)  
    { return m_map.find(key) != m_map.end(); }  
};
```

```
Map m; m.contain("a very long string"); // Temporary std::string created
```

->

```
class Map {  
    std::map<std::string, int, std::less<>> m_map; // C++14 transparent comparator  
public:  
    bool contain(std::string_view key) {  
        { return m_map.find(key) != m_map.end(); }  
    };  
Map m; m.contain("a very long string"sv);
```

Using std::string_view with std::unordered_map is possible thanks to the technique shown by Marc Mutz, "StringViews, StringViews everywhere!" at "Meeting C++ 2017". C++20 makes it easy: P0919.

Use gsl::span

```
class TwoDTable { ...  
    void lookup(double x, std::vector<double>& y) const;  
};  
std::vector<double> val(2);  
m_table->lookup(x, val);
```

->

```
class TwoDTable { ...  
    void lookup(double x, ix::span<double> y) const;  
};  
std::array<double,2> val = {};  
m_table->lookup(x, val);
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`gsl::span`, `std::string_view`).
 - **Use moves.**
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Use move

```
{  
    std::vector<std::int64_t> vertex_ids_for_this_cell;  
    ...  
    cell_vertex_ids.push_back(vertex_ids_for_this_cell); // Copy!  
}
```

->

```
{  
    std::vector<std::int64_t> vertex_ids_for_this_cell;  
    ...  
    cell_vertex_ids.push_back(std::move(vertex_ids_for_this_cell));  
}
```

Use move: rule of zero

```
class X {  
    ~X(); // Disable move operations generation  
    std::string m_s;  
};  
X::~~X() = default;  
std::vector<X> v = ...;  
std::reverse(begin(v), end(v)); // temporary copies are created  
->  
class X {  
    std::string m_s; // Rule of zero  
};  
static_assert(std::is_nothrow_move_constructible_v<X>);  
std::vector<X> v = ...;  
std::reverse(begin(v), end(v)); // Move operations are used.
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`gsl::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - **Use `std::array`**, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Use std::array

```
void calc(...) { ...  
    std::vector<double> COF(15, 0.0);  
    COF[0] = ...  
    ...  
    COF[14] = ...
```

->

```
void calc(...) {  
    std::array<double, 15> COF = {  
        ...  
    };
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`gsl::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, **`boost::container::small_vector`**.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Vocabulary types: `small_vector`

- Popularised by LLVM.
CppCon 2016: Chandler Carruth “High Performance Code 201: Hybrid Data Structures”
- Implemented as **`boost::container::small_vector`** (by Ion Gaztanaga)
- `small_vector<T, 8>`
 - Will handle any size
 - Will not allocate any memory as long as the size never exceeds 8.
 - Is clearly larger than `vector<T>`
 - Meets STL requirement even when T is “bool”.


```
template <typename T, int N>
class SmallVector
    : public SmallVectorImpl<T> {
    char Buffer[sizeof(T) * N];

public:
    SmallVector()
        : SmallVectorImpl((T *)Buffer,
                          (T *)Buffer,
                          N) {}

    // ...
};
```

```
template <typename T>
class SmallVectorImpl {
    T *Begin, *End;
    size_t Capacity;

protected:
    SmallVectorImpl(T *Begin, T *End,
                   size_t Capacity);

public:
    iterator begin() { return Begin; }
    iterator end() { return End; }

    void push_back(const T &Element);
    void pop_back();

    // ...
};
```

(6 / 40)



CHANDLER CARRUTH

High Performance Code 201: Hybrid Data Structures

5:27 / 55:48

CppCon.org

boost::container::small_vector - overhead

```
// With gcc 64 bit:
```

```
// sizeof(vector<int>)          = 24, sizeof(vector<double>)      = 24
// sizeof(small_vector<int, 1>) = 32, sizeof(small_vector<double, 1>) = 32
// sizeof(small_vector<int, 2>) = 32, sizeof(small_vector<double, 2>) = 40
// sizeof(small_vector<int, 3>) = 32, sizeof(small_vector<double, 3>) = 48
// sizeof(small_vector<int, 4>) = 40, sizeof(small_vector<double, 4>) = 56
// sizeof(small_vector<int, 5>) = 40, sizeof(small_vector<double, 5>) = 64
// sizeof(small_vector<int, 6>) = 48, sizeof(small_vector<double, 6>) = 72
// sizeof(small_vector<int, 7>) = 48, sizeof(small_vector<double, 7>) = 80
// sizeof(small_vector<int, 8>) = 56, sizeof(small_vector<double, 8>) = 88
// ...
// sizeof(small_vector<int,15>) = 80, sizeof(small_vector<double,15>) = 144
// sizeof(small_vector<int,16>) = 88, sizeof(small_vector<double,16>) = 152
```

```
// With gcc 64 bit:
```

```
// sizeof(std::vector<bool>) = 40, sizeof(std::vector<char>) = 24
// N      sizeof(small_vector<bool, N>)
//  1:15  32 bytes
// 16:23  40 bytes
// 24:31  48 bytes
```

span + small_vector

- Step 1: replace **vector** by **span** in function parameters
Any contiguous containers can now be used.
- Step 2: replace vector by **small_vector** where necessary and if beneficial

Case study – problem #2 – small_vector

170e6 allocations

```
std::vector<int> js(n);
```

->

```
boost::container::small_vector<int, 16> js(n);
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`gsl::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - **Avoid `pimpl` when necessary. Use `std::optional`.**
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Avoiding memory allocations

- Avoid Pimpl when too expensive
 - Pimpl are about data hiding
 - Not worth it for internal code that is in the critical path
- Use `std::optional` for member data if delayed initialisation is necessary

Avoid Pimpl when necessary

```
class NodeContainerIterator {  
    struct IIterator;  
    std::shared_ptr<IIterator> m_pimpl;
```

->

```
class NodeContainerIterator {  
    std::vector<Node*>::const_iterator m_iterator;
```

Use std::optional to delay initialisation

```
class ExecutorCommand {  
    bool pre_execution(...);  
    // "Context" not default initialisable. Initialised in pre_execution().  
    std::unique_ptr<Context> m_ctx;
```

->

```
class ExecutorCommand {  
    bool pre_execution(...);  
    // Avoid memory allocation  
    std::optional<Context> m_ctx;
```


Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`gsl::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`
 - Avoid `pimpl` when necessary. Use `std::optional`.
- **Re-use allocated memory.**
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Re-use vector capacity

```
for (...) {  
    std::vector<int> bsz;  
    for (...)  
        bsz.push_back(...);
```

->

```
std::vector<int> bsz; // hoisted to re-use capacity  
for (...) {  
    bsz.clear();      // clear it  
    bsz.reserve(...); // “reserve” safe after a clear()  
    for (...)  
        bsz.push_back(...);
```

The same technique can be used with `std::string` and `std::stringstream`.

Do **not** use reserve() in a loop

```
for (...) {  
    auto block = received();  
    if (block.empty())  
        break;  
buffer.reserve(buffer.size() + block.size());  
    for (auto const& elt : block)  
        buffer.push_back(elt);
```



Performance killer

Create expensive objects only once

```
std::ostringstream oss;  
oss.imbue(std::locale(oss.getLoc(), new boost::posix_time::time_facet("%d-%b-%Y")));
```

->

```
using tf = boost::posix_time::time_facet;  
static const auto owned_time_facet =  
    std::make_unique<tf>  
    ("%d-%b-%Y",  
     tf::period_formatter_type{}, tf::special_values_formatter_type{}, tf::date_gen_formatter_type{},  
     1); // Non zero ref count to retain ownership  
std::ostringstream oss;  
oss.imbue(std::locale(oss.getLoc(), owned_time_facet.get()));
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`gsl::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Avoid node-based containers

```
void process(std::vector<int> const& recvIndexSet)
{
```

```
    std::unordered_set<int> receivedNodes;
```

Flagged as a source
of allocations

```
    for (int i : recvIndexSet) {
```

```
        // There is a chance of duplication in recvIndexSet.
```

```
        if (receivedNodes.insert(i).second) {
```

```
            ...
```

Avoid node-based containers

```
std::sort(begin(recv_ind), end(recv_ind));  
process(recv_ind);
```



Call site

```
void process(std::vector<int> const& recvIndexSet)  
{  
    std::set<int> receivedNodes;  
    for (int i : recvIndexSet) {  
        // There is a chance of duplication in recvIndexSet.  
        if (receivedNodes.find(i) == receivedNodes.end()) {  
            receivedNodes.insert(i);  
            ...  
        }  
    }  
}
```

Avoid node-based containers

```
// Sort and remove duplicates
sort(begin(recv_ind), end(recv_ind));
recv_ind.erase(std::unique(begin(recv_ind), end(recv_ind)),
                end(recv_ind));
process(recv_ind);

void process(std::vector<int> const& recvIndexSet)
// Expect: recvIndexSet is sorted and contains no duplicates
{
    ...
}
```


Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`gsl::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Part III: C++17 “pmr” allocators

- C++17 introduces the “polymorphic memory resources”
- Available in VS 2017 and GCC 9.
- Nicolai Josuttis “C++17 - The complete Guide”, chapter 31

Howard Hinnant’s `stack_alloc` is still handy.

(https://howardhinnant.github.io/stack_alloc.html)

C++17 pmr allocators - vector

// Allocated and never resized

```
std::vector<double> f1(num_1);
```

```
std::vector<double> f2(num_1);
```

```
std::vector<double> f3(num_1*num_1);
```

```
std::vector<double> f4(num_2);
```

```
std::vector<double> f5(num_2*num_2);
```

->

```
std::array<char, 256> stack_buffer;
```

```
std::pmr::monotonic_buffer_resource mem_res{std::data(stack_buffer), std::size(stack_buffer)};
```

```
std::pmr::vector<double> f1(num_1, 0., &mem_res);
```

```
std::pmr::vector<double> f2(num_1, 0., &mem_res);
```

```
std::pmr::vector<double> f3(num_1*num_1, 0., &mem_res);
```

```
std::pmr::vector<double> f4(num_2, 0., &mem_res);
```

```
std::pmr::vector<double> f5(num_2*num_2, 0., &mem_res);
```

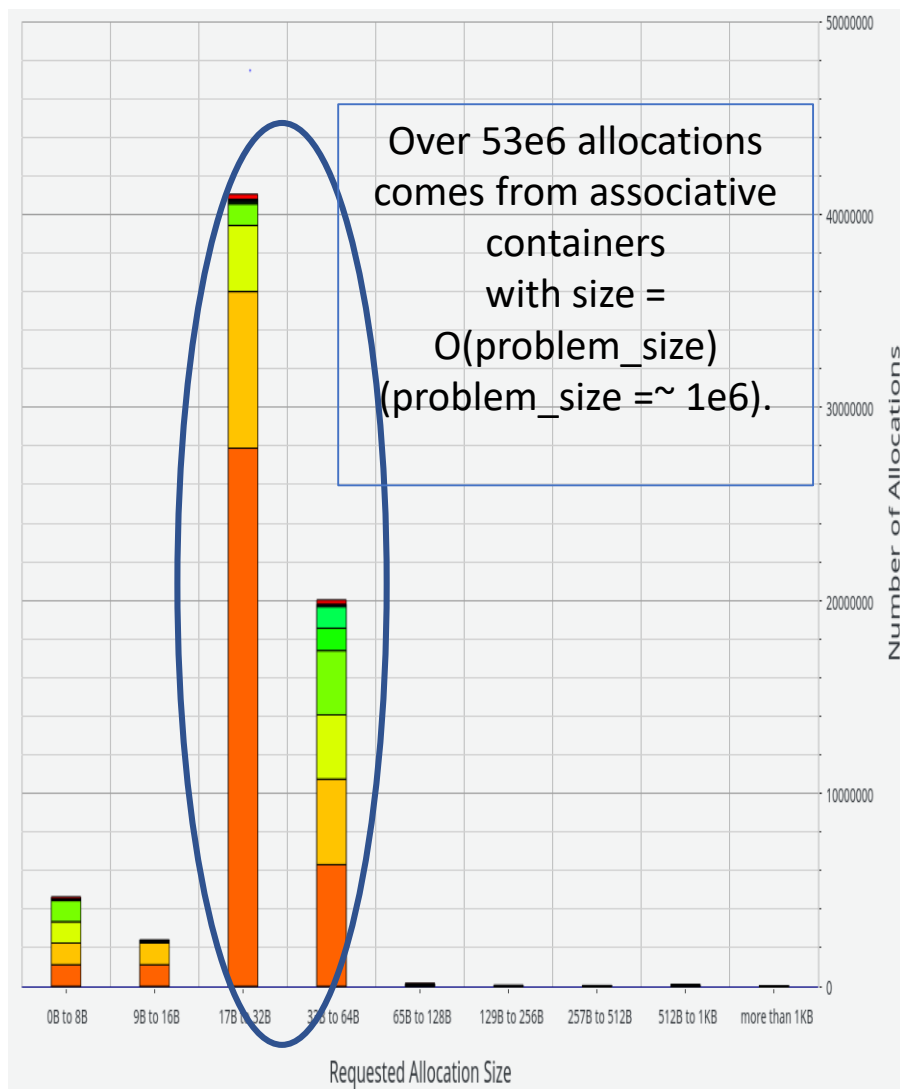
C++17 pmr allocators - map

```
class X {  
    // No individual erase operations  
    std::unordered_map<std::int64_t, std::size_t> m_id_to_index;  
public:  
    X() {}  
};  
->
```

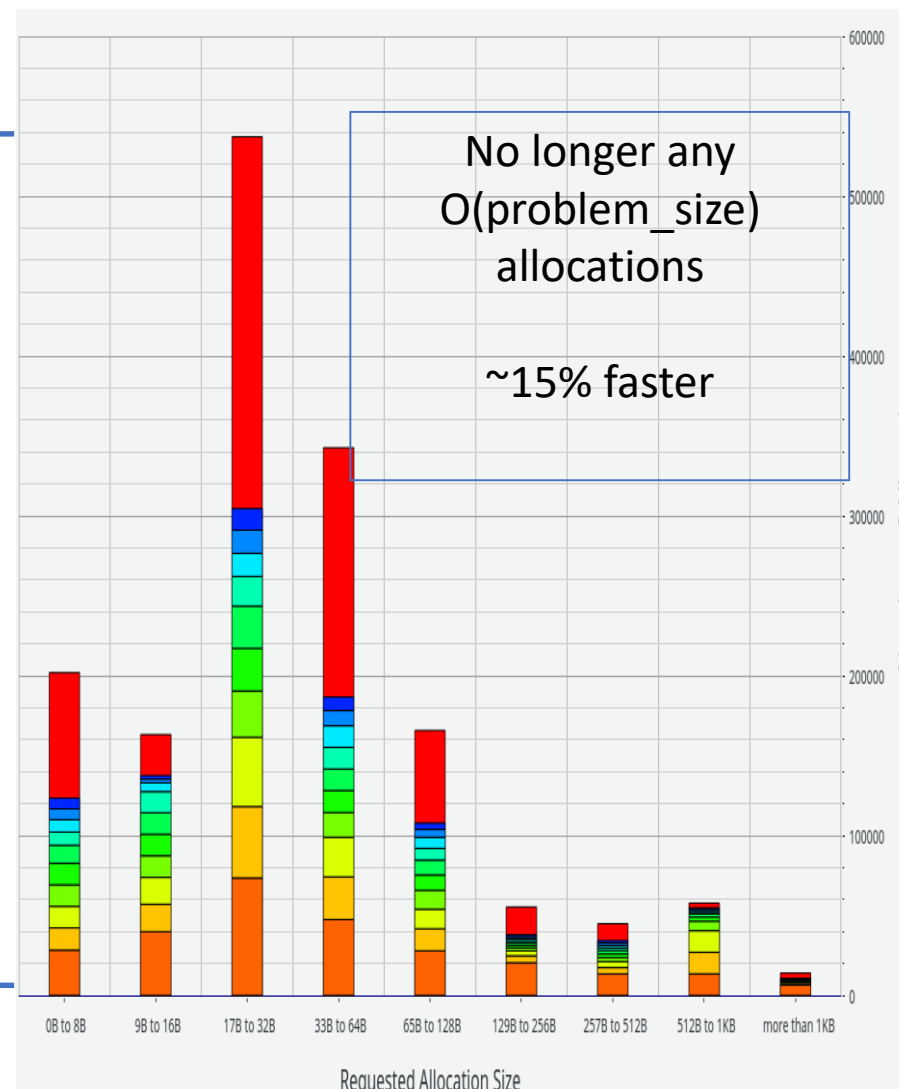
```
class X {  
    std::pmr::monotonic_buffer_resource m_res;  
    std::pmr::unordered_map<std::int64_t, std::size_t> m_id_to_index;  
public:  
    X() : m_id_to_index{&m_res} {}  
};
```

~50e6

~0.6e6



Before



After

Anti-pattern: monotonic_buffer_resource and emplace/insert

```
class Cache {  
    std::pmr::monotonic_buffer_resource m_res;  
    std::pmr::unordered_map<std::size_t, int> m_cache;  
public:  
    X() : m_cache{&m_res} {}  
    int get_index(std::size_t key) {  
        // A temporary node is created resulting in slow memory growth.  
        const auto r = m_cache.emplace(key, -1);  
        if (!r.second)  
            r.first->second = get_new_index(...);  
        return r.second;  
    }  
};
```



Memory killer

Pattern: monotonic_buffer_resource and try_emplace

```
class Cache {
    std::pmr::monotonic_buffer_resource m_res;
    std::pmr::unordered_map<std::size_t, int> m_cache;
public:
    X() : m_cache{&m_res} {}
    int get_index(std::size_t key) {
        // C++17 try_emplace does not create any temporary node.
        const auto r = m_cache.try_emplace(key, -1);
        if (!r.second)
            r.first->second = get_new_index(...);
        return r.second;
    }
};
```

Conclusions

- Memory allocations can be surprisingly expensive.
- Heaptrack rocks.
- Do not copy unnecessarily.
- Learn about views (`string_view`, `span`) and `small_vector`.

Appendix

C++17 `std::string_view`

Supported by VS 2017, gcc 7.1 and later

For older compilers:

- boost's "[boost/utility/string_view.hpp](https://github.com/boostorg/string_view)"
- https://github.com/tcbrindle/cpp17_headers
- `std::string` compatible `std::hash` support missing but easy to implement

`gsl::span`, C++20 `std::span`

- <https://github.com/microsoft/GSL>
- <https://github.com/tcbrindle/span>
- <https://github.com/martinmoene/span-lite>