

Reducing Memory Allocations In A Large C++ Application

Arnaud Desitter

CppOnSea Conference

15 July 2020

In memoriam

Hubert Matthews

RIP 2019



Bio

Arnaud Desitter

arnaud.desitter@gmail.com

Lives in Oxford, UK

25 years experience in working on scientific software

Developer for Schlumberger on a reservoir simulator

Opinions in this talk are my own, not my employer's.

Long standing ACCU member

And a cyclist !

Roadmap

- Motivations
- Methodology
- Part I: a case study
- Part II: solutions to address excessive memory allocations
 - Vocabulary types
 - Patterns and anti-patterns
- Part III: C++17 pmr allocators (briefly)
- Conclusions

Why optimise memory allocations ?

- Improve **performance** (usually speed)
- Decrease **memory footprint**

Performance may be of little importance.

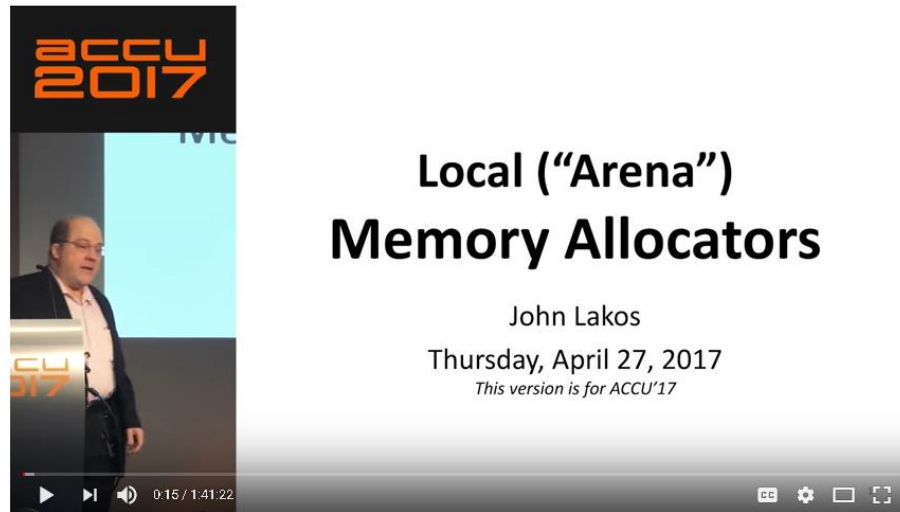
Performance may not be a priority.

There may be better ways to improve performance:
find hot spot, use better algorithms, etc.

That said, large applications tend to have a rather flat profile and perform too many memory allocations.

Motivations

Custom allocators are a much discussed topic in the C++ industry.




Local (arena) Memory Allocators - John Lakos [ACCU 2017]

Motivations

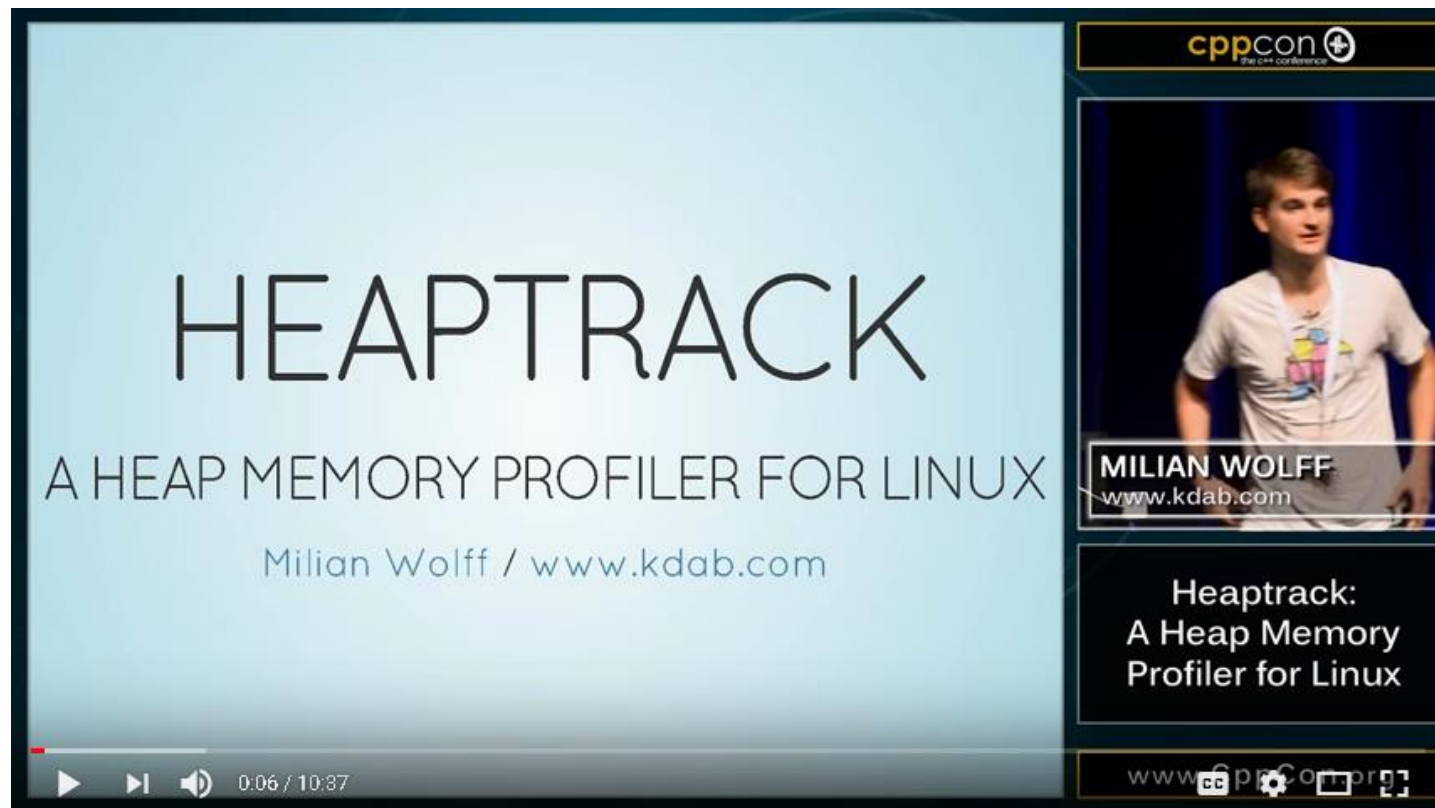
Custom allocators are a much discussed topic in the C++ industry.

How do I quantify the memory allocations of my application ?



Local (arena) Memory Allocators - John Lakos [ACCU 2017]

Motivations



CppCon 2015: Milian Wolff "Heaptrack: A Heap Memory Profiler for Linux"

Installing Milian Wolff's heaptrack

- Build it from source
search for “heaptrack build ubuntu”
- Use pre-built AppImage
 - https://download.kde.org/stable/heaptrack/1.1.0/heaptrack-v1.1.0-x86_64.AppImage.mirrorlist
 - <https://travis-ci.org/KDAB/heaptrack> for recent snapshots

```
wget .../heaptrack-v1.1.0-x86_64.AppImage  
chmod +x heaptrack-v1.1.0-x86_64.AppImage
```

Running Milian Wolff's heaptrack

heaptrack operates on any executable.

Stack traces available only with debugging symbols (compile with “-g”)

```
> ./heaptrack-v1.1.0-x86_64.AppImage /bin/ls
```

heaptrack output will be written to "heaptrack.ls.1877.zst"

starting application, this might take some time...

heaptrack.ls.1877.zst **heaptrack-v1.1.0-x86_64.AppImage**

heaptrack stats:

allocations: 44

leaked allocations: 38

temporary allocations: 2

Heaptrack finished! Now run the following to investigate the data:

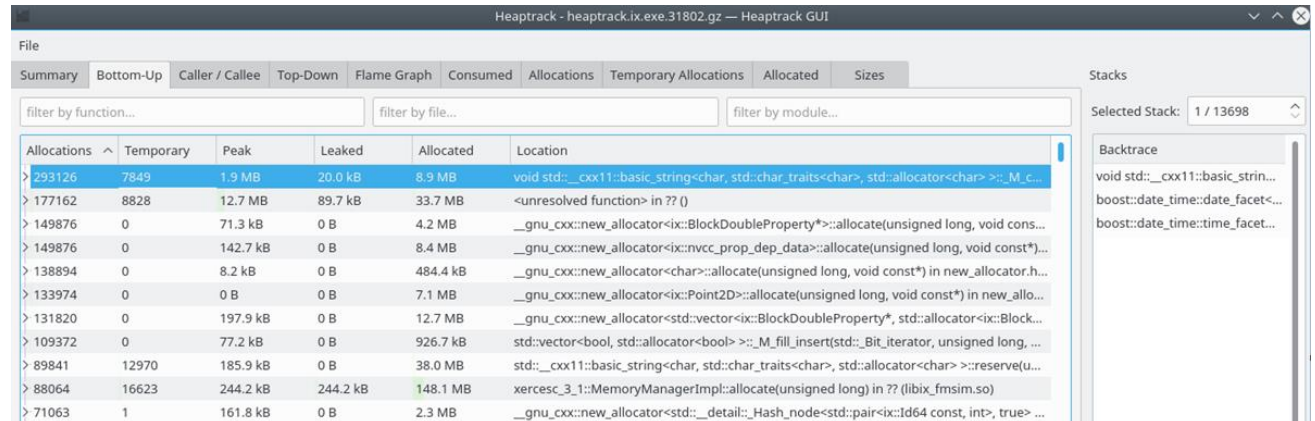
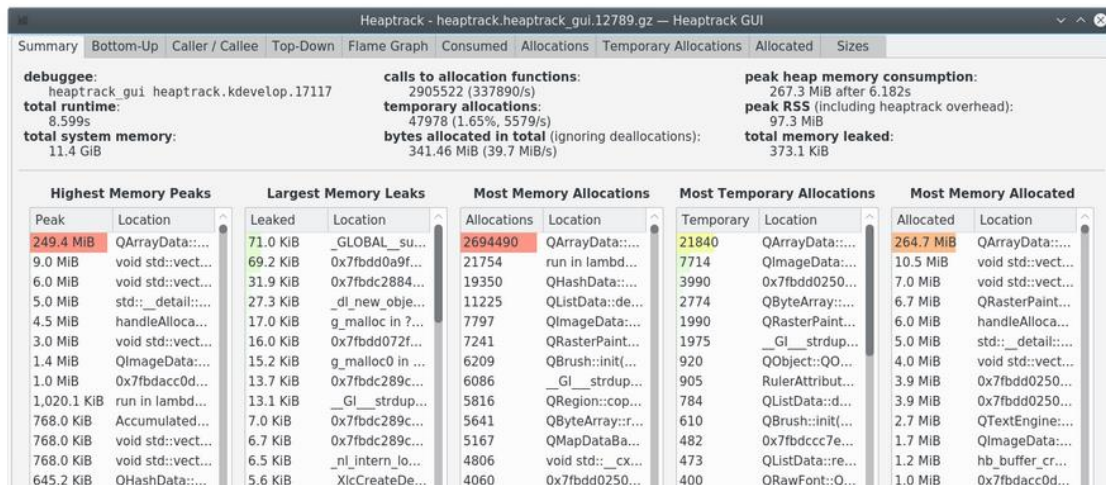
```
heaptrack --analyze "heaptrack.ls.1877.zst"
```

```
> ./heaptrack-v1.1.0-x86_64.AppImage --analyze heaptrack.ls.1877.zst
```

(1) Execution with
heaptrack data collection

(2) Data
visualisation


Heaptrack




Heaptrack




Flamecharts



Cumulated allocations



Sizes



Consumed

Demo time

Profiling: a methodology

- Choose a case.
- Run it under a profiler.
- Spot a problem.
- Try to fix it.
- Profile again with the fix.
 - Discard fix if it does not work.
 - Submit if it does.
- Iterate until there are no more opportunities for change.
- Choose another case and iterate.

Part I: a case study


A numerical simulation

- Reasonably large synthetic problem
... but shortened to a single day of simulation
- Dominated by floating point computation
- Run **without** any concurrency for the sake of this example

Reservoir simulator

- Considered amongst the most advanced of its type
- Constantly developed for more than 15 years
- Multi-millions lines C++ code base
- No memory leaks

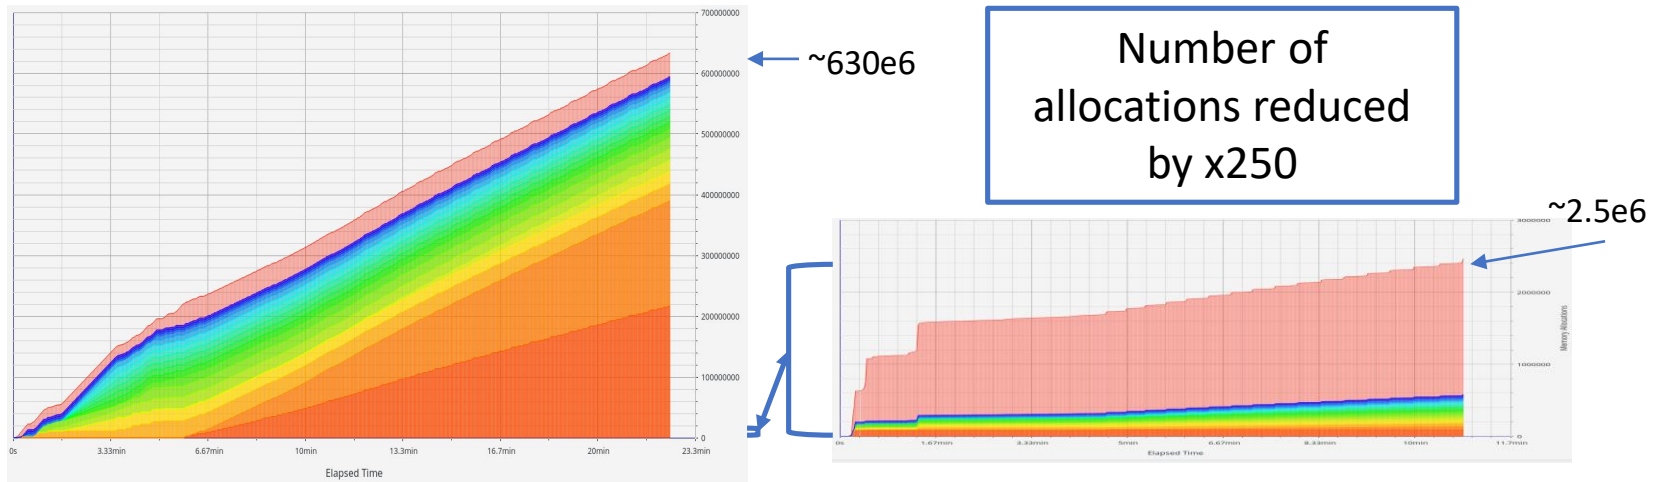
A case study



← ~630e6

Cumulative
number of
allocations

A case study




Before

After

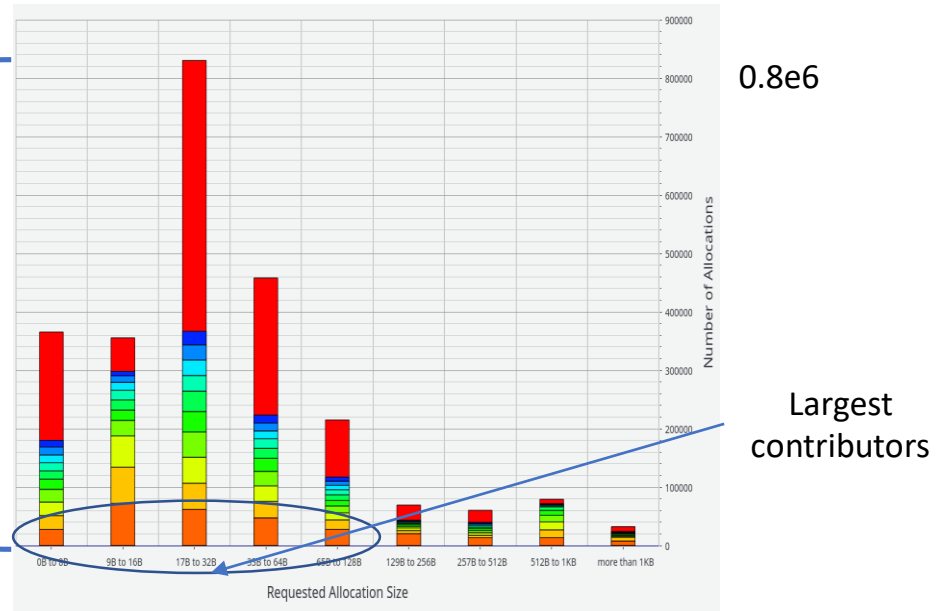
A case study

Most allocations were for 8 bytes or less.

Map nodes



Before




After

Costs of allocations are hard to predict

- Before: **5.5%** time was spent in malloc/free.
- After: **1.5%**.
- But we have a speed up well above 10%.

Reported time spent in allocator is an **under-estimate** of possible gain, as excessive memory allocations can be a source of cache-misses.

Continuous performance improvement



2018 delivers continuous improvement in performance by better memory management and optimization in the linear solver. Benchmarking over a wide range of field models confirm an average performance improvement of 15% over 2017 release.

Memory growth: a case study

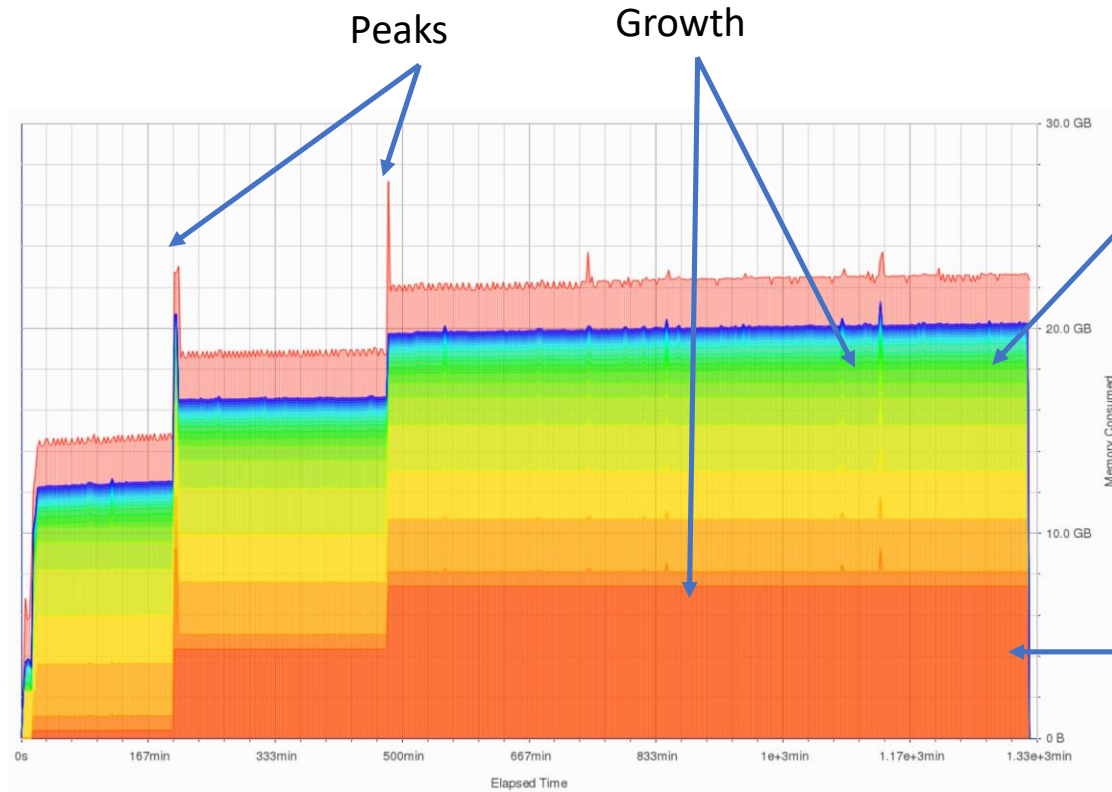
A large numerical simulation

- Real world case. Very large grid.
- 128 processes. About 15 Gb per process.
- ~18 h

As the current version of the simulator calls the memory allocator infrequently, using heaptrack is possible.

Memory growth: a case study - before


~ 18 h, 128
processes
using about
15 Gb each



Due to
retaining
vector
capacity for
too long

Due to a very
pessimistic
vector<>::reserve()
on a long living
object

Memory growth: a case study - after



Conclusions so far

- Heaptrack works really well.
 - Uses **debug information** provided by gcc and clang (“-g”).
 - **Small overhead** once the number of allocations is under control.
 - Compatible with **threads** and inter-process communication.
- It helped us to identify and then obtain significant speedup
 - Biggest gains come from **removing allocations within loops**.
 - Biased towards small size allocations.

Part II: solutions

- Do not copy if you can.
- Avoid allocations.
- Re-use allocated memory.
- Use contiguous containers.

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Part II: solutions

- Do not copy if you can.
 - **Avoid unused objects.**
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Case study – problem #1 – Avoid unnecessary copies

Some leftovers of a debugging experiment: 210e6 allocations

```
std::vector<double> average_m(m.size());  
for (int ic = 0; ic < m.size(); ic++)  
    average_m[ic] = m[ic];
```



```
auto const& average_m = m;
```

ALWAYS DO LESS WORK



EFFICIENCY WITH ALGORITHMS, PERFORMANCE WITH DATA STRUCTURES

Chandler Carruth

▶ ▶▶ 🔊 32:16 / 1:13:40



CppCon 2014: Chandler Carruth "Efficiency with Algorithms, Performance with Data Structures"

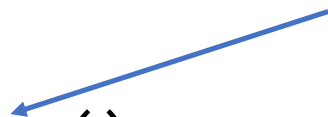
Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - **Use references.**
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Do not copy – “auto const& x = f()”

```
for (...) { for (...) {  
    std::vector<int> bs =  
    the_sizes[...]->get_sizes();
```

Returns “std::vector<int> const&”



```
for (...) { for (...) {  
    // See GotW #88: A Candidate For the “Most Important const”  
    auto const& bs = the_sizes[...]->get_sizes();
```

Do not copy arguments

```
// copy !
```

```
void add_res_from_buffer(std::vector<double> dbuffer);
```



```
void add_res_from_buffer  
    (std::vector<double> const& dbuffer);
```


Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - **Use views (`tcb::span`, `std::string_view`).**
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Vocabulary types - views

- Views are non-owning **pointers**.
 - They can dangle.
 - They are cheap to copy.
- C++17 `std::string_view`
 - A read-only view to a **contiguous** sequence of “char”.
 - May **not** have a trailing ‘\0’.
- C++20 `std::span<T>`
 - Implemented by Tristan Brindle as `tcb::span<T>`
 - A view to a **contiguous** sequence of T.
 - `tcb::span<T>`: read-write
 - `tcb::span<const T>`: read-only

Vocabulary types – `std::string_view`

- Avoid allocating a string when referring to a substring
- Avoid creating temporary `std::string` when passing a literal (“A very long string”).
“`std::string const &`” -> “`std::string_view`” can be beneficial.

Use std::string_view

```
int size = 0;  
const std::unique_ptr<char[]> pchar = block->get_char(size);  
std::string str(pchar.get(), size - 1); // Copy!
```



```
int size = 0;  
const std::unique_ptr<char[]> pchar = block->get_char(size);  
const std::string_view str{ pchar.get(), size - 1 };
```

string_view and std::map

```
struct Map {  
    std::map<std::string, int> m_map;  
    bool contain(std::string const& key) const  
    { return m_map.find(key) != m_map.end(); }  
};  
Map m; m.contain("a very long string");
```

Temporary
std::string
created

```
struct Map {  
    std::map<std::string, int, std::less<>> m_map;  
    bool contain(std::string_view key) const  
    { return m_map.find(key) != m_map.end(); }  
};  
Map m; m.contain("a very long string"sv);
```

C++14
transparent
comparator

string_view and std::unordered_map

C++20 makes it easy: P0919.

Possible in C++17 using:

Marc Mutz, *StringViews, StringViews everywhere!*, Meeting C++ 2017

```
std::unordered_map<std::string, int> map;
```



```
std::unordered_map<  
    std::string_view, std::pair<std::string, int>> map;
```




Vocabulary types – tcb::span

- Usually used as function parameters
 - “std::vector<double> **&**”
-> tcb::span<double> (if std::vector is not resized)
 - “std::vector<double> **const&**”
-> tcb::span<**const** double>
- Allow passing any contiguous sequences such as:
 - std::vector,
 - std::array,
 - boost small_vector, static_vector,
 - a sub-section of a contiguous sequence,
 - etc.

Use tcb::span

```
class TwoDTable { ...  
    void lookup(double x, std::vector<double>& y) const;  
};  
std::vector<double> val(2);  
m_table->lookup(x, val);
```



```
class TwoDTable { ...  
    void lookup(double x, tcb::span<double> y) const;  
};  
std::array<double,2> val = {};  
m_table->lookup(x, val);
```


Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - **Use moves.**
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Use “move”

```
{  
    std::vector<std::int64_t> vertex_ids_for_this_cell;  
    ...  
    cell_vertex_ids.push_back(vertex_ids_for_this_cell); // Copy!  
}
```



```
{  
    std::vector<std::int64_t> vertex_ids_for_this_cell;  
    ...  
    cell_vertex_ids.push_back(std::move(vertex_ids_for_this_cell));  
}
```

Use “move”: rule of zero

```
class X {  
    ~X() = default; // Disable move operations generation  
    std::string m_s;  
};  
std::vector<X> v = ...;  
std::reverse(begin(v), end(v)); // temporary copies are created
```



```
class X {  
    std::string m_s; // Rule of zero  
};  
static_assert(std::is_nothrow_move_constructible_v<X>);  
std::vector<X> v = ...;  
std::reverse(begin(v), end(v)); // Move operations are used.
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - **Use `std::array`**, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Use std::array

```
void calc(...) { ...  
    std::vector<double> COF(15, 0.0);  
    COF[0] = ...  
    ...  
    COF[14] = ...
```



```
void calc(...) {  
    std::array<double, 15> COF = {  
        ...  
    };
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, **`boost::container::small_vector`**.
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Vocabulary types: `small_vector`

- Implements the **small buffer optimisation** using the `std::vector` API
- `small_vector<T, N>`
 - Will handle any size.
 - Will not allocate any memory as long as the size never exceeds `N`.
 - Is larger than `vector<T>`.
- Popularised by LLVM.
CppCon 2016: Chandler Carruth “High Performance Code 201: Hybrid Data Structures”

```
template <typename T, int N>
class SmallVector
    : public SmallVectorImpl<T> {
    char Buffer[sizeof(T) * N];

public:
    SmallVector()
        : SmallVectorImpl((T *)Buffer,
                          (T *)Buffer,
                          N) {}

    // ...
};
```

```
template <typename T>
class SmallVectorImpl {
    T *Begin, *End;
    size_t Capacity;

protected:
    SmallVectorImpl(T *Begin, T *End,
                   size_t Capacity);

public:
    iterator begin() { return Begin; }
    iterator end() { return End; }

    void push_back(const T &Element);
    void pop_back();

    // ...
};
```

(6 / 40)



CHANDLER CARRUTH

High Performance Code 201: Hybrid Data Structures

5:27 / 55:48

CppCon.com

CppCon 2016: Chandler Carruth "High Performance Code 201: Hybrid Data Structures"

Boost small_vector

Implemented as **boost::container::small_vector**
(by Ion Gaztanaga)

Header only, exception safe, no specialisation for bool

```
// Propagation of noexcept requires boost 1.71 or later
#define BOOST_MOVE_HAS_NO_THROW_MOVE_ASSIGN(T) \
    std::is_nothrow_move_assignable_v<T>
#include <boost/container/small_vector.hpp>
```

Case study – problem #2 – small_vector

170e6 allocations

```
std::vector<int> js(n);
```



```
namespace bc = boost::container;  
bc::small_vector<int, 16> js(n);
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - **Avoid pimpl when necessary.** Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.


Avoid Pimpl when too expensive

- Pimpl are about data hiding.
- Not worth it for internal code that is in the critical path.

Avoid Pimpl when necessary

```
class NodeContainerIterator {  
    struct IIterator;  
    std::shared_ptr<IIterator> m_pimpl;
```

Flagged as an
excessive source
of allocations



```
class NodeContainerIterator {  
    std::vector<Node*>::const_iterator  
    m_iterator;
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`.
 - Avoid `pimpl` when necessary. **Use `std::optional`.**
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Use std::optional to delay initialisation

```
class ExecutorCommand {  
    // “Context” not default initialisable.  
    // Initialised in pre_execution().  
    std::unique_ptr<Context> m_ctx;  
    bool pre_execution(...);  
}
```



```
class ExecutorCommand {  
    std::optional<Context> m_ctx; // Avoid memory allocation  
    bool pre_execution(...);  
}
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`
 - Avoid `pimpl` when necessary. Use `std::optional`.
- **Re-use allocated memory.**
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Re-use vector capacity

```
for (...) {  
    std::vector<int> bsz;  
    for (...)  
        bsz.push_back(...);  
}
```




```
std::vector<int> bsz; // hoisted to re-use capacity  
for (...) {  
    bsz.clear();      // clear it  
    bsz.reserve(...); // “reserve” safe after a clear()  
    for (...)  
        bsz.push_back(...);  
}
```

The same technique can be used with `std::string` and `std::stringstream`.

Do **not** use `reserve()` in a loop

```
for (...) {  
    auto const& blocks = received();  
    if (blocks.empty())  
        break;  
    buffer.reserve(buffer.size() +  
        blocks.size());  
    for (auto const& block : blocks)  
        buffer.push_back(block);  
}
```

Performance killer



Replace push_backs by insert

```
for (...) {  
    auto const& blocks = received();  
    if (blocks.empty())  
        break;  
for (auto const& block : blocks)  
    buffer.push_back(block);  
    buffer.insert(end(buffer),  
        begin(blocks), end(blocks));  
}
```


Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Avoid node-based containers

```
void process(std::vector<int> const&
            recvIndexSet)
{
    std::unordered_set<int> receivedNodes;
    for (int i : recvIndexSet) {
        // Avoid duplicates in recvIndexSet
        if (receivedNodes.insert(i).second) {
            ...
        }
    }
}
```

Flagged as a source of
allocations



Avoid node-based containers

```
std::sort(begin(recv_ind), end(recv_ind));  
process(recv_ind);
```

Call site

```
void process(std::vector<int> const&  
            recvIndexSet)  
{  
    ...  
}
```

Avoid node-based containers

```
// Sort and remove duplicates
```

```
std::sort(begin(recv_ind), end(recv_ind));  
recv_ind.erase(std::unique(begin(recv_ind),  
                             end(recv_ind)),  
                end(recv_ind));  
process(recv_ind);
```

```
void process(std::vector<int> const& recvIndexSet)  
// Expect: recvIndexSet is sorted and contains no  
//         duplicates  
{
```

Part II: solutions

- Do not copy if you can.
 - Avoid unused objects.
 - Use references.
 - Use views (`tcb::span`, `std::string_view`).
 - Use moves.
- Avoid allocation.
 - Use `std::array`, `boost::container::small_vector`
 - Avoid `pimpl` when necessary. Use `std::optional`.
- Re-use allocated memory.
 - Use `std::vector::reserve()`.
 - Make use of `std::vector` capacity.
- Use contiguous containers.
 - Avoid if possible `std::map`, `std::set` and `std::list` in critical code.

Part III: C++17 “pmr” allocators

C++17 introduces the “polymorphic memory resources”.

- Available in VS 2017 and GCC 9
(be aware of <https://gcc.gnu.org/PR94906>, fixed in gcc 9.4)
- Nicely explained in:
Nicolai Josuttis, *C++17 - The Complete Guide*, chapter 29

Howard Hinnant’s `stack_alloc` is still handy.

(https://howardhinnant.github.io/stack_alloc.html)

C++17 “pmr” allocators

- pmr and vector
 - A use case
- pmr and node based containers
 - A use case
 - An anti-pattern

C++17 pmr allocators - vector

// Allocated and never resized

```
std::vector<double>  
  f1(num_1), f2(num_1*num_1),  
  f3(num_2), f4(num_2*num_2);
```



```
std::array<std::byte, 256> stack_buffer;  
std::pmr::monotonic_buffer_resource mem_res  
  {std::data(stack_buffer), std::size(stack_buffer)};  
std::pmr::vector<double>  
  f1(num_1, 0., &mem_res), f2(num_1*num_1, 0., &mem_res),  
  f3(num_2, 0., &mem_res), f4(num_2*num_2, 0., &mem_res);
```

C++17 pmr allocators - map

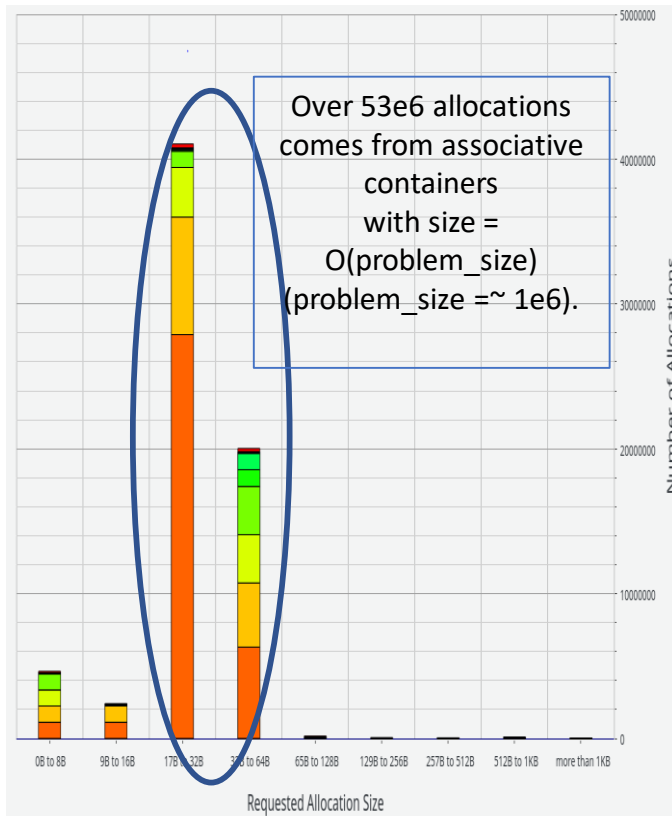
```
class X {  
    // No individual erase operations  
    std::unordered_map<std::int64_t, std::size_t> m_id_to_index;  
public:  
    X() = default;  
};
```



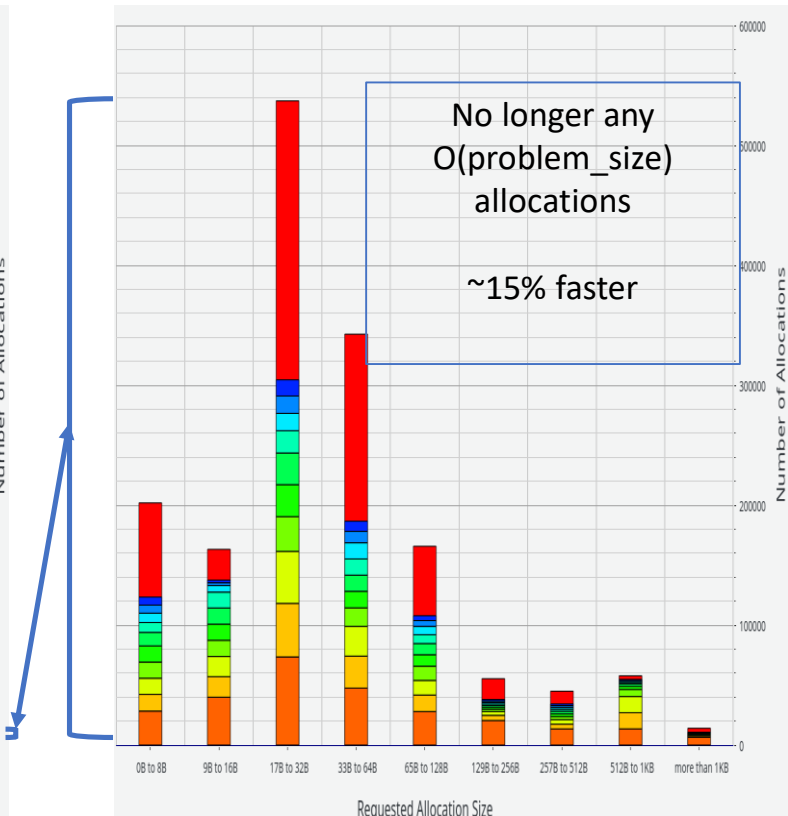
```
class X {  
    std::pmr::monotonic_buffer_resource m_res;  
    std::pmr::unordered_map<std::int64_t, std::size_t> m_id_to_index{&m_res};  
public:  
    X() = default;  
    X(X const&) = delete; X& operator=(X const&) = delete;  
};
```

~50e6

~0.6e6



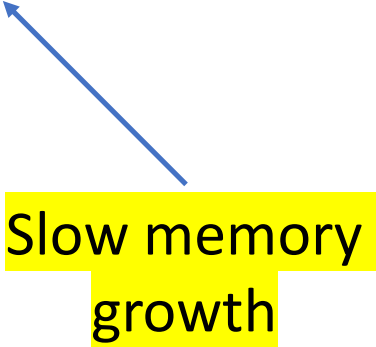
Before



After

Anti-pattern: monotonic_buffer_resource and emplace/insert

```
class Cache {  
    std::pmr::monotonic_buffer_resource m_res;  
    std::pmr::unordered_map<std::size_t, int> m_cache{&m_res};  
public:  
    int get_index(std::size_t key) {  
        // A temporary node is created resulting in slow memory growth.  
        const auto [it, inserted] = m_cache.emplace(key, -1);  
        if (!inserted)  
            it->second = get_new_index(...);  
        return it->second;  
    }  
};
```



Slow memory growth

Pattern: monotonic_buffer_resource and try_emplace

```
class Cache {
    std::pmr::monotonic_buffer_resource m_res;
    std::pmr::unordered_map<std::size_t, int> m_cache{&m_res};
public:
    int get_index(std::size_t key) {
        // C++17 try_emplace does not create any temporary node.
        const auto [it, inserted] = m_cache.try_emplace(key, -1);
        if (!inserted)
            it->second = get_new_index(...);
        return it->second;
    }
};
```

Conclusions

- Memory allocations can be surprisingly expensive.
- Heaptrack rocks.
- Top tips
 - Do not copy unnecessarily.
 - Learn about views (`string_view`, `span`) and `small_vector`.

Appendix

C++17 `std::string_view`

Supported by VS 2017, gcc 7.1 and later

For older compilers:

- boost's "`boost/utility/string_view.hpp`"
- https://github.com/tcbrindle/cpp17_headers
- `std::string` compatible `std::hash` support missing but easy to implement

`tcb::span`, `gsl::span`, C++20 `std::span`

- <https://github.com/tcbrindle/span>
- <https://github.com/microsoft/GSL>
- <https://github.com/martinmoene/span-lite>